

RecSplit

Minimal Perfect Hashing

A new algorithm that is faster and uses less space

Contents

- Use Cases
- Hashing a Fixed Set
- Universal Hashing
- (Minimal) Perfect Hashing
- RecSplit
- BDZ, CHD

Use Cases

- You have a large, fixed (immutable) set of keys
- Keys don't fit in memory (many millions)
- Values are small (e.g. 8 bits/entry, or possibly just one bit if you do garbage collection)
- You need a small in-memory index (e.g. 2 bits/key)

Use Cases: Examples

- Page Rank: billions of URLs, each with an 8 bit counter, and need keep it in memory for speed
- Garbage collection in a de-duplication store: billion of keys, each with one bit, to mark used entries
- Eye-based tracking: each word has many thousand "near matches" (tracking is very inaccurate)

Possible Solutions

- Hashing can have collisions, so we can't use it
- Perfect hashing: maps n keys, without collision, to an index $0..m$, where $m > n$ (gaps are possible)
- Minimal perfect hashing: map n keys, without conflicts and without gaps, to a index $0..n$
- k -perfect hashing: each key has at most k collisions

Hashing a Fixed Set

		Collisions	Gaps
RecSplit {	Regular Hashing	Yes	Yes
	Perfect Hashing	No	Yes
	Minimal Perfect Hashing	No	No
	k -Perfect Hashing	at most k	No

Perfect Hashing

Disadvantages:

- Only works for static sets
- Requires $O(n)$ memory for the hash function description
- Keys are not stored; can not check if a key is in the set (if false positives are OK, can store a hash "fingerprint")

Universal Hashing

- To generate a perfect hash function (PHF), can not rely on a simple hash function, due to possible conflicts
- Universal hashing can create endless many hash codes for each key, using an indexed (seeded) hash function
- All (minimal-) perfect hashing algorithms use universal hashing

Universal Hashing

Regular hash function (Java):

```
public static int hashCode(byte a[]) {  
    int result = 1;  
    for (byte element : a)  
        result = 31 * result + element;  
    return result;  
}
```

Universal Hashing

Universal hash function (simplified):

```
public static int universalHash(byte a[], int index) {  
    int result = index;  
    for (byte element : a)  
        result = 31 * result + element + index;  
    return result;  
}
```

- This sample implementation is not very "random"
- Better use more secure algorithms, such as seeded Murmur Hash, SipHash, SHA-256,...

Universal Hashing

- Example: list of $\text{universalHash}(\text{key}, \text{index}) \bmod 4$

Index:	0	1	2	3	4	5	6	7
Key: "a"	1	2	3	2	3	2	3	2
Key: "b"	0	1	2	2	2	2	2	1
Key: "c"	0	3	2	2	3	1	1	0
Key: "d"	2	1	2	2	0	1	0	0

Universal Hashing

- $f(k) = \text{universalHash}(k, 6) \bmod 4$
is a MPHf (minimal perfect hash function)

Index:	0	1	2	3	4	5	6	7
Key: "a"	1	2	3	2	3	2	3	2
Key: "b"	0	1	2	2	2	2	2	1
Key: "c"	0	3	2	2	3	1	1	0
Key: "d"	2	1	2	2	0	1	0	0

- Brute force algorithm: find the first index where $\text{universalHash}(\text{key}, \text{index}) \bmod 4$ has no collision

(Minimal) Perfect Hashing

- So generating an MPH is simple for small sets: try `universalHash(x)`, and if there is a collision, try with $x+1, \dots$ until there is no collision
- But for larger sets, too many tries are needed (millions for sets of size > 20)
- The challenge is to support huge sets, $O(n)$ generation, and constant time evaluation

Best Algorithms

- Theoretical lower bound: 1.44 bits/key
- RecSplit: around 1.8 bits/key
- CHD: around 2.1 bits/key
- BDZ: around 2.7 bits/key

(assuming somewhat reasonable generation time; both RecSplit and CHD can approach the theoretical lower bound given enough time)

RecSplit

The algorithm has three phases:

- Partitioning
- Bucket Processing
- Storing

RecSplit Partitioning

- A. Partition the n keys into $n/100$ buckets
- B. Overflow: buckets with more than 2000 keys are merged and processed with BDZ (fallback). This is an extremely rare case, and makes RecSplit a *pseudo* hybrid algorithm: less than 0.0000...001% of the keys fall into this category (in practise: none).
- C. What remains is bucket with up to 2000 keys

RecSplit Partitioning

Keys

jan	feb	mar	apr	may	jun	jul	aug	sep	oct	nov	dec
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----

RecSplit Partitioning

Keys

jan	feb	mar	apr	may	jun	jul	aug	sep	oct	nov	dec
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----



$\text{universalHash}(k, 0) \bmod \text{bucketCount}$



Buckets

0

1

2

3

RecSplit Partitioning

Keys

jan	feb	mar	apr	may	jun	jul	aug	sep	oct	nov	dec
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----



$\text{universalHash}(k, 0) \bmod \text{bucketCount}$

Buckets

0

feb	mar
may	jun
jul	dec

1

jan	nov
-----	-----

2

apr	aug
sep	oct

3

(empty)

RecSplit Partitioning

Keys

jan	feb	mar	apr	may	jun	jul	aug	sep	oct	nov	dec
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----



$\text{universalHash}(k, 0) \bmod \text{bucketCount}$

Buckets

0

1

2

3

feb	mar
may	jun
jul	dec

jan	nov
-----	-----

apr	aug
sep	oct

(empty)

Overflow

(empty, as no bucket is overly large)

RecSplit Bucket

Convert each bucket into a tree recursively:

- A. Sets of size 0 and 1: no processing is needed
- B. Sets up to size 8 (or so): use the "brute force" algorithm described in Universal Hashing
- C. Larger (size s): find the first x so that $\text{universalHash}(\text{key}, x)$ splits the sets in two (sizes $s/2$ and rest), and process each set recursively

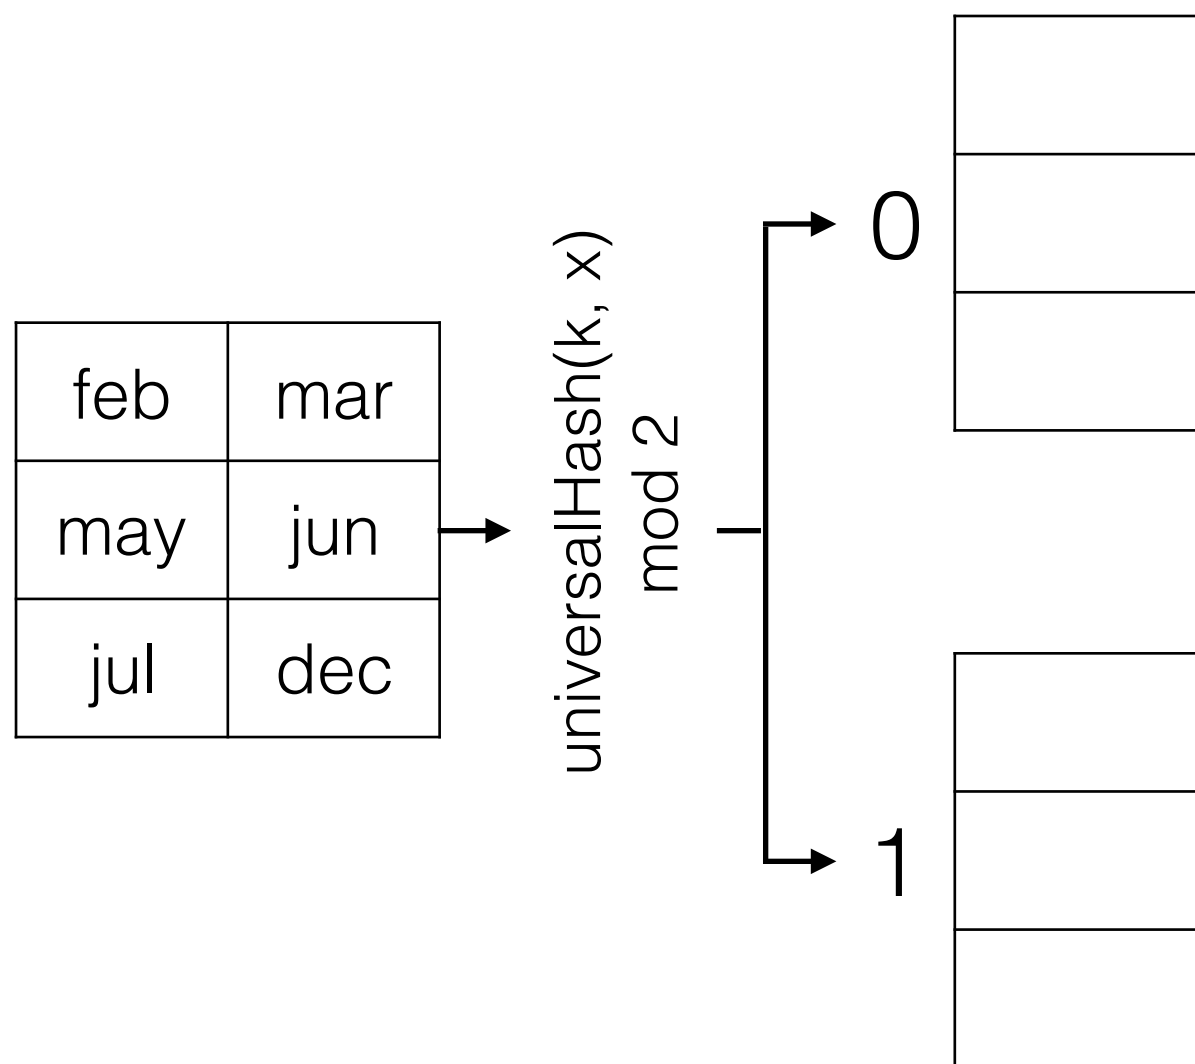
RecSplit Bucket

Bucket

feb	mar
may	jun
jul	dec

RecSplit Bucket

Bucket



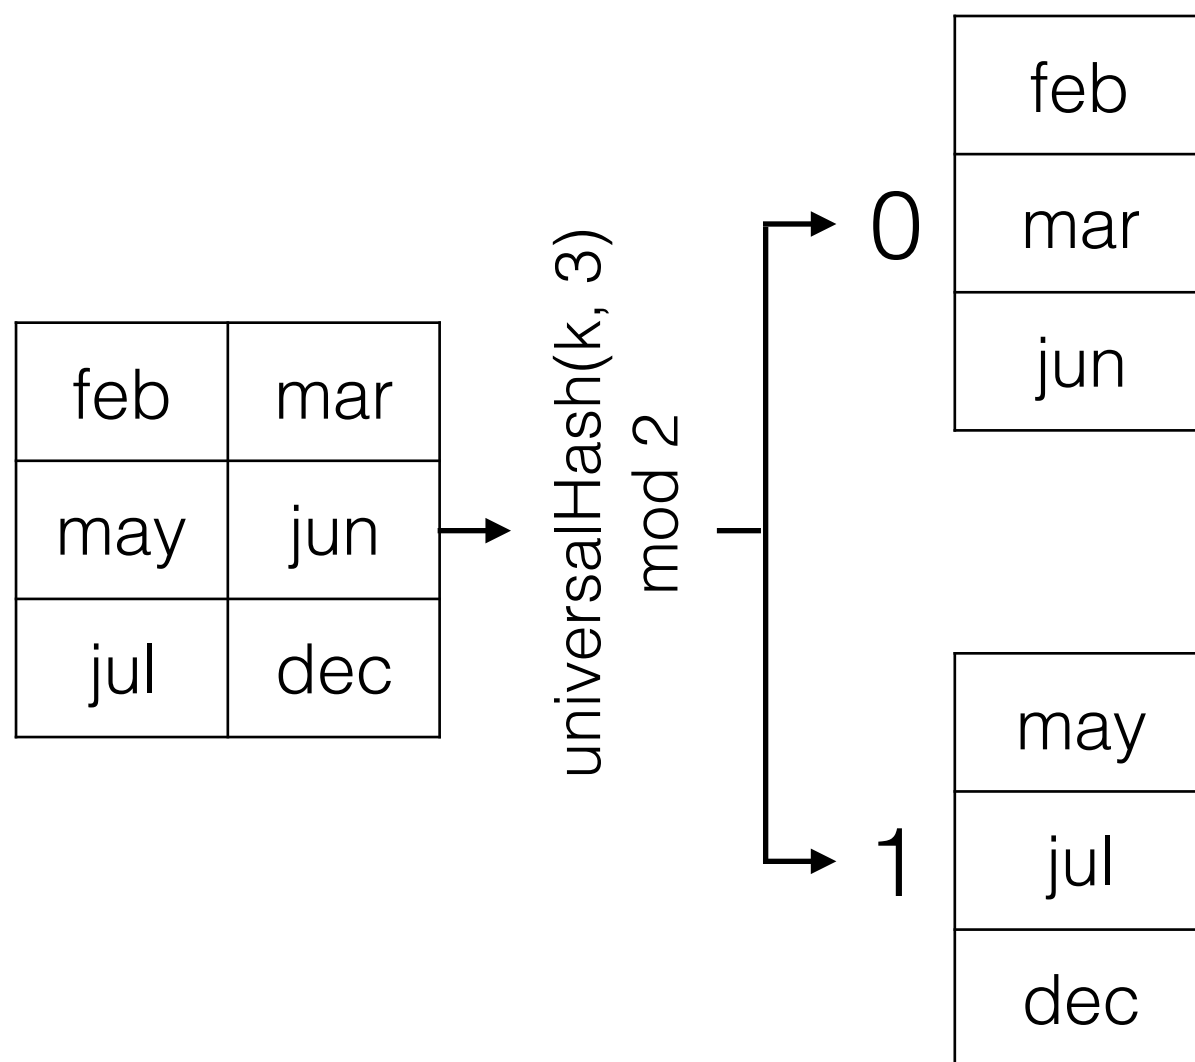
Search the first x that evenly splits the set

(for an odd size, the first subset has one entry more)

RecSplit Bucket

Bucket

Bucket Description: 3,

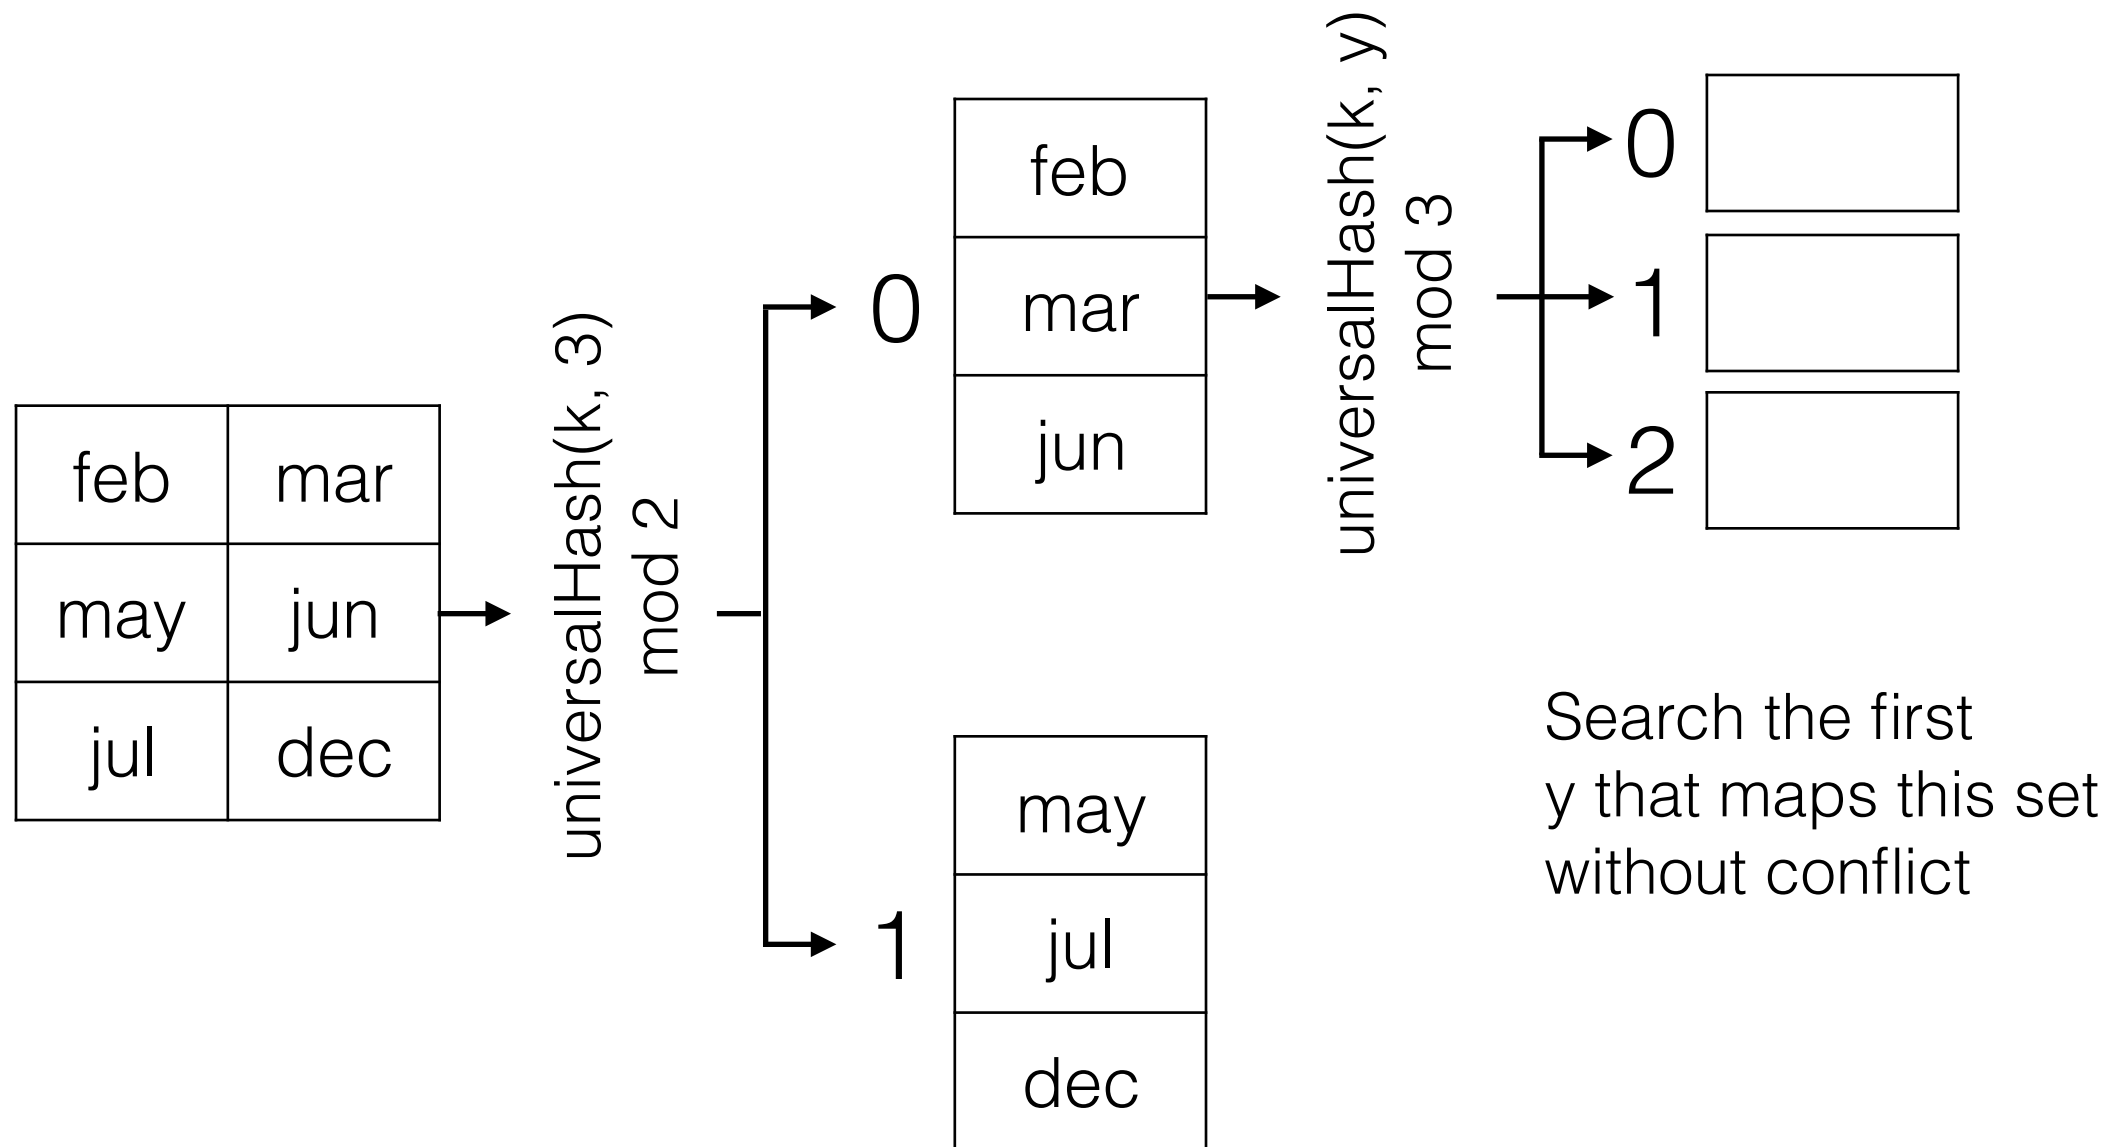


We find that $\text{universalHash}(k, 3)$ splits the set as needed, so 3 is the first entry in the bucket description

RecSplit Bucket

Bucket

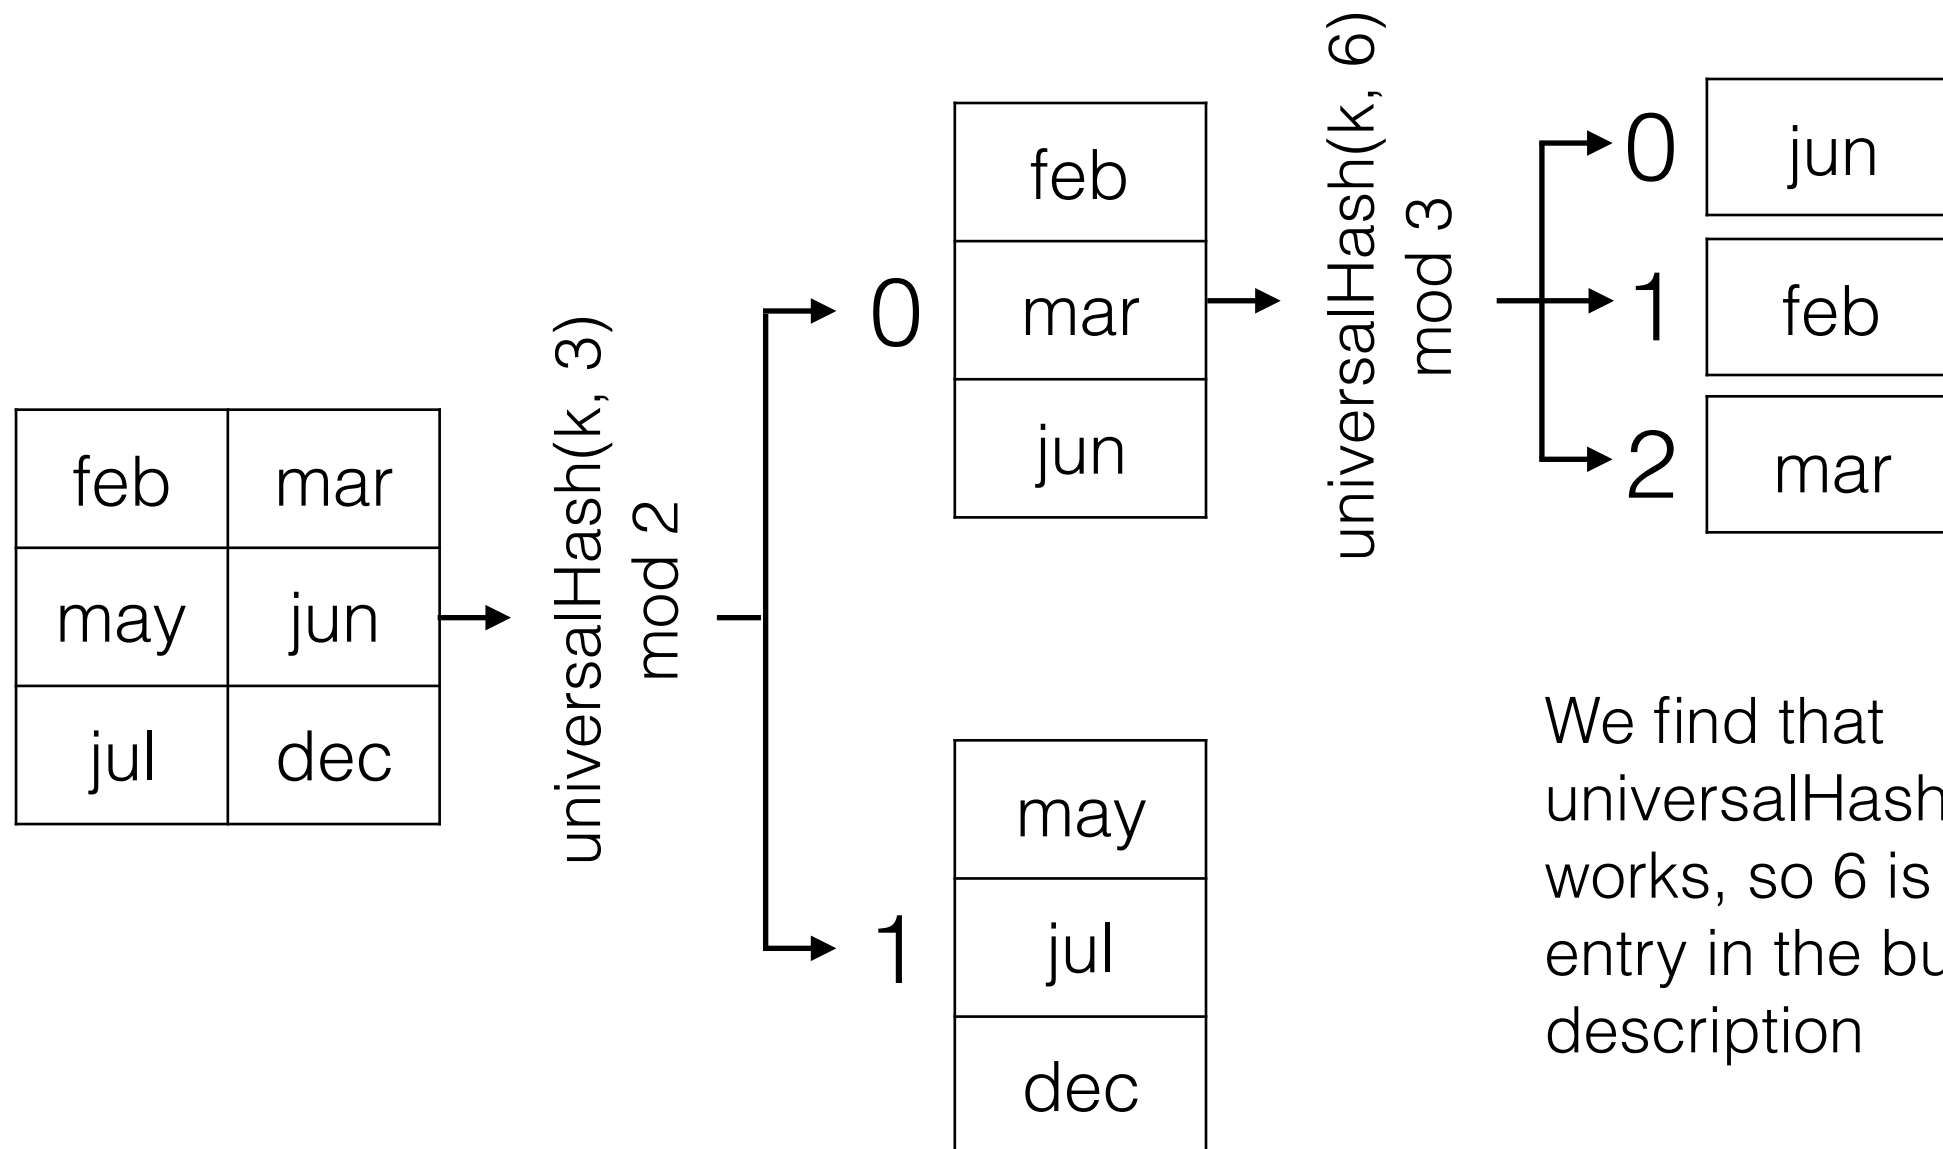
Bucket Description: 3,



RecSplit Bucket

Bucket

Bucket Description: 3, 6,

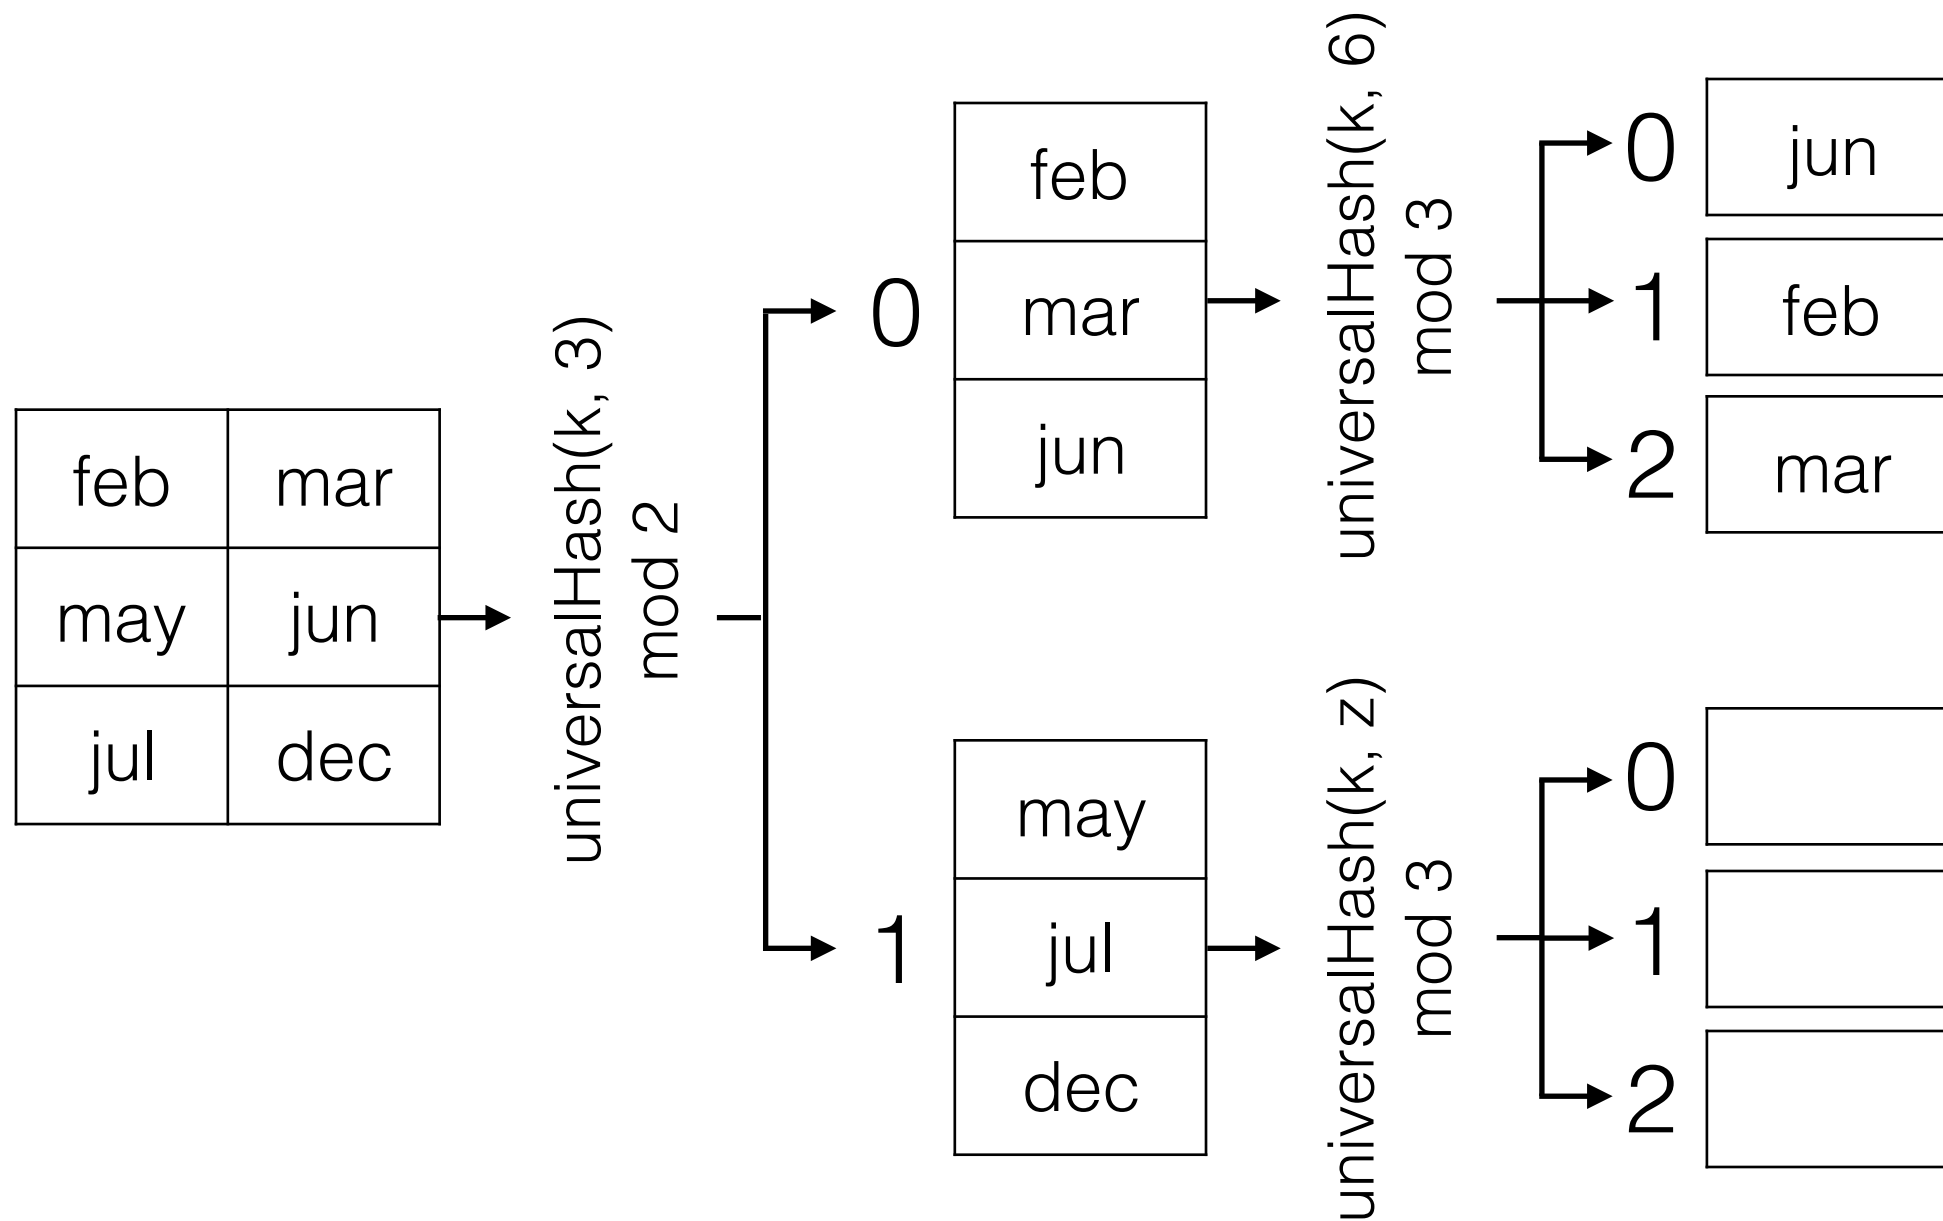


We find that $\text{universalHash}(k, 6)$ works, so 6 is the next entry in the bucket description

RecSplit Bucket

Bucket

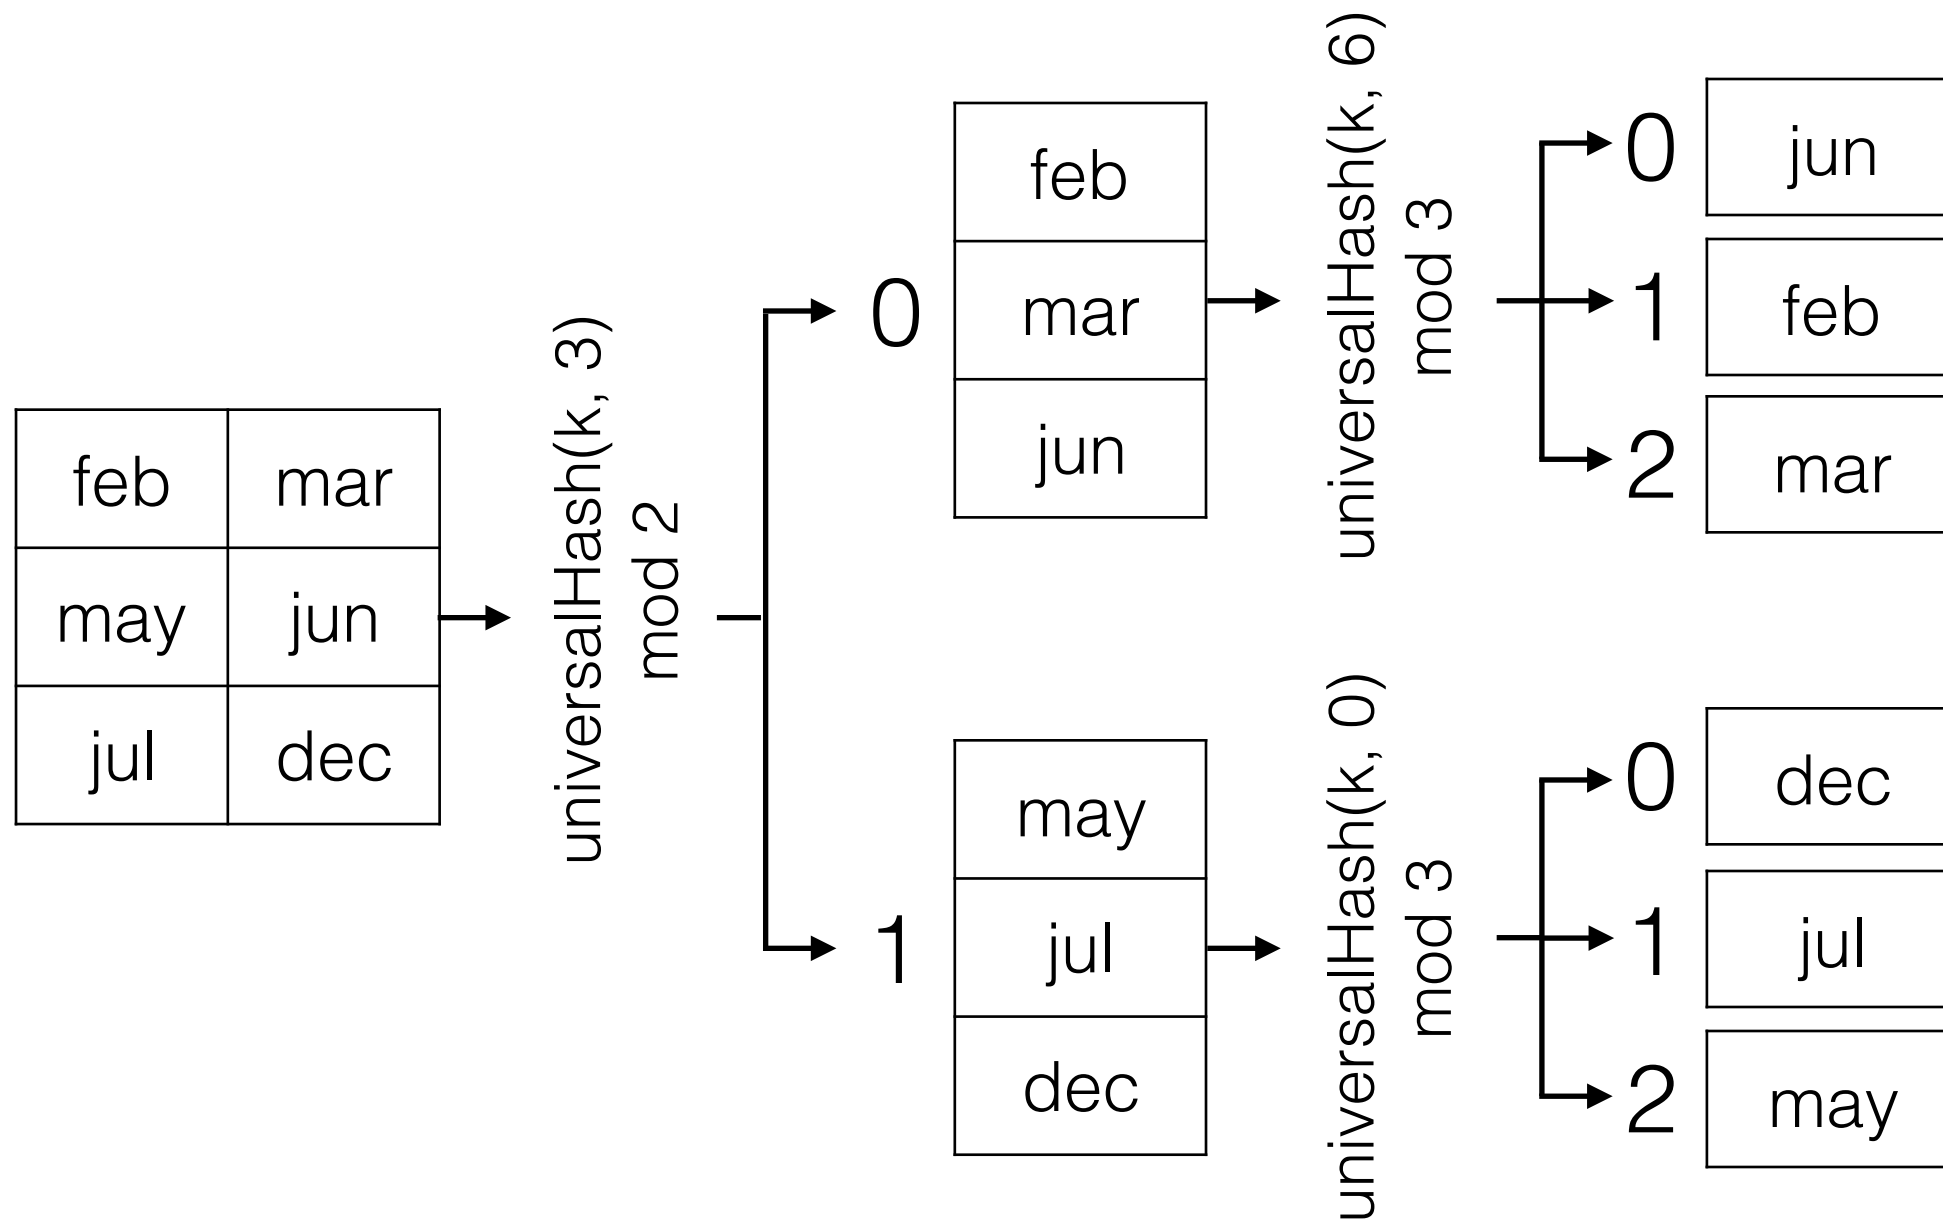
Bucket Description: 3, 6,



RecSplit Bucket

Bucket

Bucket Description: 3, 6, 0

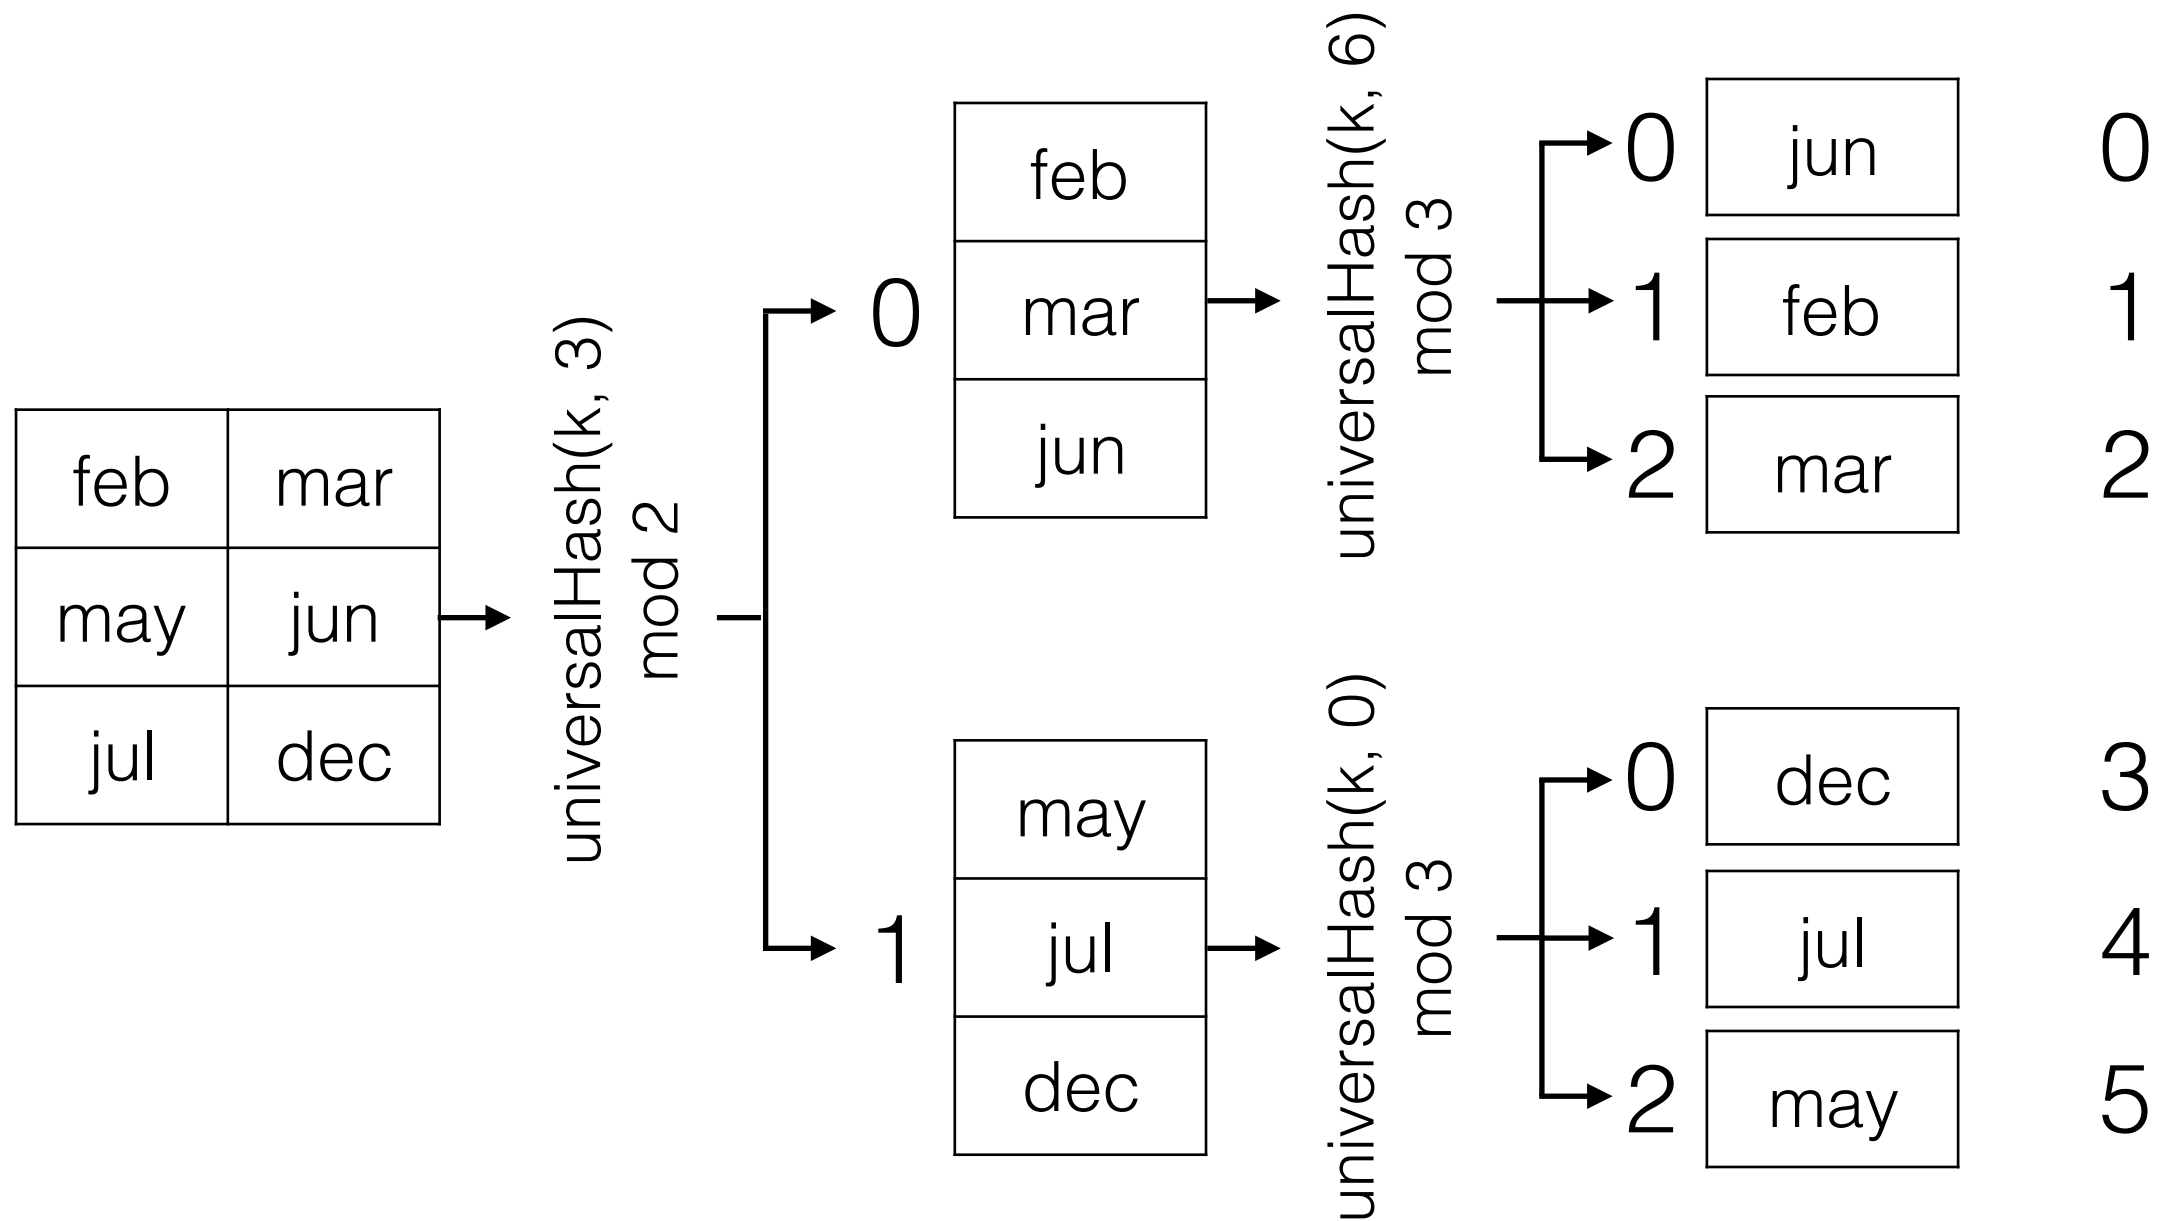


RecSplit Bucket

Bucket

Bucket Description: 3, 6, 0

ID



RecSplit Storing

- A. Store the set size (using the Elias Delta code)
- B. Store the bucket lookup table and bucket size table (as Elias-Fano monotone lists)
- C. Store each bucket description: A bucket description is the list of universal hash indexes that describe how to split the tree (stored using the Rice code)

RecSplit Observations

- RecSplit is easy to parallelise as buckets are processed independently
- The average bucket size, and how to process buckets, is configurable, which allows to trade space for generation time and evaluation speed

RecSplit (k -)Perfect

RecSplit is a MPHF (minimal perfect hash function), but small modifications allow for:

- (non-minimal) perfect hash, by changing the last stage to use a larger modulo
- k -perfect hash, by changing the last stage to use a smaller modulo and allow for k entries per slot

BDZ

<http://cmph.sourceforge.net/bdz.html>

- A. For n keys, prepare $1.23n$ slots
- B. For each key, calculate 3 potential slots h_0, h_1, h_2
- C. For each slot that contains only one key,
remove the key from the *other* slot that also contain it,
leaving just one key per slot
- D. This will probably work; if not, repeat at B with
a new universalHash index
- E. For each slot, store whether it contains h_0, h_1 , or h_2
- F. Use a Rank data structure to map the $1.23n$ slots to $0..n$

BDZ Observations

- Related to Cockoo Hashing
- Hard to parallelise because slots need to be processed in order
- If it doesn't work ($p < 0.5$), start from scratch
- First generate a perfect hash function (PHF), and then map that to a minimum perfect hash function (MPHF) using a rank data structure
- Needs much more space than RecSplit

CHD

<http://cmph.sourceforge.net/chd.html>

- A. Partition the n keys into $n/4$ buckets
- B. Prepare $1.01n$ slots
- C. For each bucket, starting with the largest, try mapping its keys to empty slots using $\text{universalHash}(k, x) \bmod \text{slotCount}$
- D. If any slot is occupied, repeat at C
- E. For each bucket, store the universalHash x used
- F. Use a Rank data structure to map the $1.01n$ slots to $0..n$

CHD Observations

- Hard to parallelise because buckets are not independent, and need to be processed in order
- Probabilistic (try-until-it-works), but usually no need to start from scratch as with BDZ
- Generate a PHF, then map to MPHF (as BDZ)
- Still needs more space than RecSplit